

Hierarchical Agglomerative Clustering

Hierarchical clustering is based on the dendrogram. In agglomerative approach, we start with each observation (data) in singleton cluster, and then sequentially merging cluster by some “linkage” methods in bottom-up fashions.

We show the two linkage methods, which are single link and complete link. For single linkage, two clusters in each step are combined together whose two **closest** members have the **smallest** “distance”. For complete linkage, two clusters in each step with the **smallest maximum** pairwise “distance” are merged.

To measure the “distance”, **datasets are converted to relational data** before clustering. We implemented four distance measurements, which are **Euclidean distance**, **Manhattan distance**, **cosine**, and **correlation**. The first two and the last two are methods for dissimilarity and similarity measurements respectively. As for cosine and correlation measures, **we do a mapping relation from similarity to dissimilarity (i.e., one minus similarity result)**.

We take Iris dataset for clustering, and the feature space is shown in Figure 1. One can find that there is only one class linearly separable from other two classes, whereas the others are not linearly separable from each other.

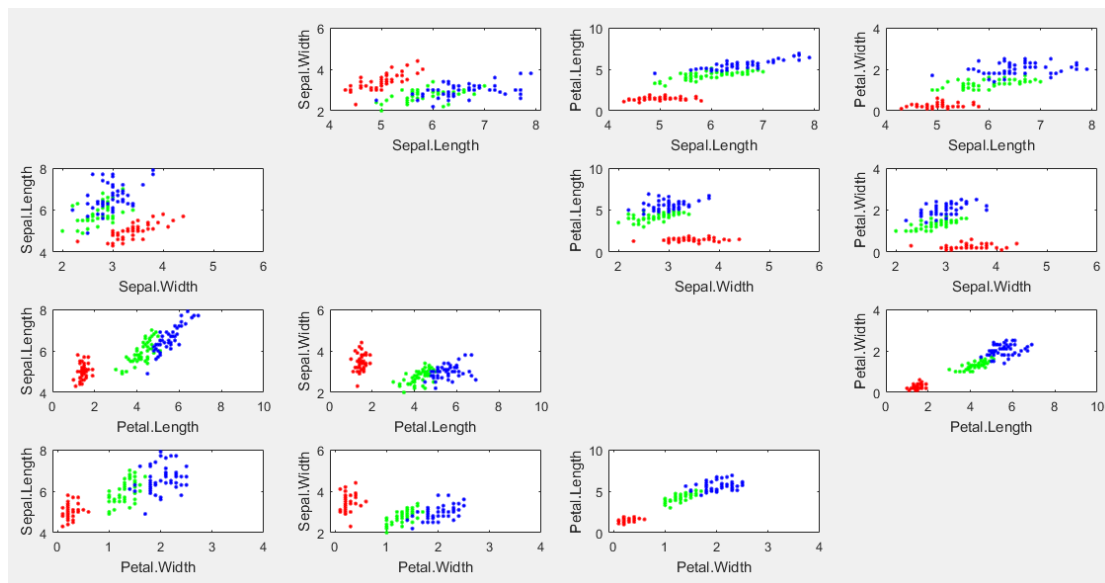


Figure 1. The feature space of Iris dataset.

In terms of linkage, **if there is a cut between the second-from-last and third-from-last linkages in one dendrogram**, then we can divide the dataset into three clusters.

Applying single-linkage clustering to Iris dataset, it tends to form two clusters as depicted in Figure 2 (i.e., one of the three clusters will include few data points). From Figure 3, we can see that data points lying in the lower right diagonal area are almost clustered into the same group. In each step for single-linkage clustering, merging two clusters with highest similarity will encourage “chaining effect”. However, similarity is usually not transitive. That is to say, if A is similar to B and B is similar to C, then it does not imply that A is similar to C. For complete-linkage clustering, it tends to break large clusters. Therefore, it works far worse than single-linkage when the data points we mentioned above belong exactly to the identical class.

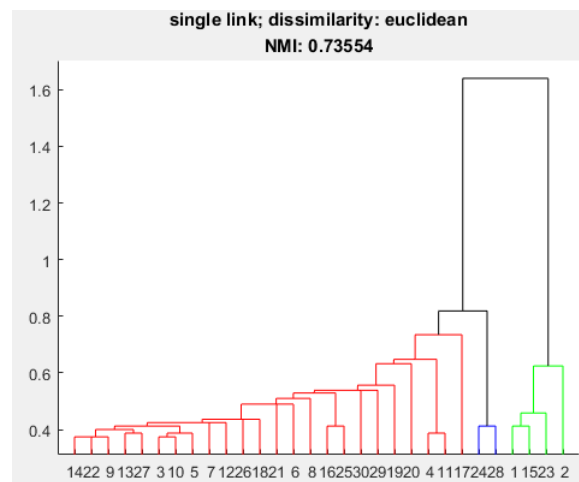


Figure 2. The dendrogram for single-linkage clustering with Euclidean measurement.

Regarding the distance measures, it makes an influence on the clustering results. A feature vector is consisted of two components, “magnitude” and “orientation”. Euclidean distance is the shortest path between two data points while Manhattan distance is the path length along each feature. Thus, they are judgments for “magnitude”, independent of orientation. For data with zero mean, cosine similarity and correlation test are exactly the same measurement. They are perceptible to “orientation”.

Table 1. Correlation coefficients of Iris dataset.

	Attribute 1	Attribute 2	Attribute 3	Attribute 4
Attribute 1	1	-0.1094	0.8718	0.8180
Attribute 2	-0.1094	1	-0.4205	-0.3565
Attribute 3	0.8718	-0.4205	1	0.9628
Attribute 4	0.8180	-0.3565	0.9628	1

As shown the correlation coefficients of each two attributes listed in Table 1, some components are in high correlation. That is to say, if two data points belong to the same

cluster, then their feature vectors might close to each other (i.e., with lower dissimilarity). However, for Euclidean distance, different vector magnitudes cause higher dissimilarity compared to that generated from correlation measurement. From Figure 3 and Figure 4, clustering results with cosine and correlation measurements are a little bit better than those adopting Euclidean distance. Notably, for Iris dataset, Manhattan distance performed the best clustering results among the four proximity measures we adopted. It can be said that this approach takes more information about the dissimilarity between two data points along each feature. Generally speaking, for high dimension data, Manhattan distance works better than Euclidean distance. Note that data points before clustering are normalized, while those shown in clustering results are actual data points.




Figure 3. The clustering results for single link.




Figure 4. The clustering results for complete link.

Interestingly, for single-linkage clustering, we try to partition Iris dataset into two clusters through a horizontal cut between the last and second-from-last linkages in the dendrogram. NMI for the clustering results are higher than those shown in Figure 3. We presented these NMIs in Table 2.

Table 2. NMIs for 2-clusters with respect to different distance measurements.

	Euclidean	Manhattan	Cosine	Correlation
Single-link	0.76117	0.76117	0.76117	0.76117

In this section, our implementation includes four proximity measures (Euclidean distance, Manhattan distance, Cosine similarity, and Correlation test), single-link clustering, complete-link clustering, and linkage for dendrogram. We use existing Matlab functions for NMI evaluation and dendrogram visualization. For either single-linkage clustering or complete linkage clustering, they all perform inferior results for the dataset we used in the previous assignments (NMIs are around 0.01 to 0.1).

DBSCAN

For DBSCAN, we first need to **determine** two parameters, *Radius* and *Pthres*. Radius is defined as the maximum distance between two samples for one to be considered as in the neighborhood of the other. Pthres is defined as the number of samples (including the point itself) in a neighborhood for a point to be considered as a core point. **The point reachable from one core point is clustered into the same group and the point assigned to no cluster is labeled as the noise point.**

Distance measurement makes great influence on the choice for *Radius*. **Assume that *Pthres* is held fixed, it is still hard to find the suitable value for *Radius*.** For convenience, we take only the mean of the distance matrix (denoted as M) as basis, and test the all selected multipliers. In our experiments, we adopted values from $0.01M$ to M with the step size $0.01M$. **In terms of *Pthres*, the larger value is, the more the number of cluster will be generated.** It is because the larger *Pthres* makes it hard to find more points in the same clusters. In our experiments, we used a **naïve way to go through all possible values for *Pthres*** (i.e., from 1 to the number of total data points).

As shown in Figure 5, we take Iris dataset for clustering and DBSCAN performs acceptable clustering results through the four proximity measures. **Except for the complete-linkage clustering result with Manhattan distance measurement, DBSCAN outperforms single-linkage and complete linkage clusterings.** Note that data points before clustering are normalized, while those shown in clustering results are actual data points **(noise points are colored in red and labeled as 0).**

Since there are **too much possible combinations** for two clustering parameters and we only simulated for some values, we **cannot give the concrete performance comparisons between these four distance measurements.** Furthermore, DBSCAN is **sensitive to clustering parameters.** We take one of the dataset we used in the previous assignments, Wireless Indoor Localization dataset (4-class 7-dimension dataset), for clustering. As shown in Figure 6, a little change of *Radius* ($0.01M$) causes larger decrease for NMI (0.10806) and so does *Pthres*. **Therefore, to choose the appropriate parameters for DBSCAN, sufficient understanding of the dataset must be required.**

All in all, compared to those agglomerative clustering approach we mentioned above, **DBSCAN shows the flexibility of clustering shapes and performs relatively acceptable clustering results.**

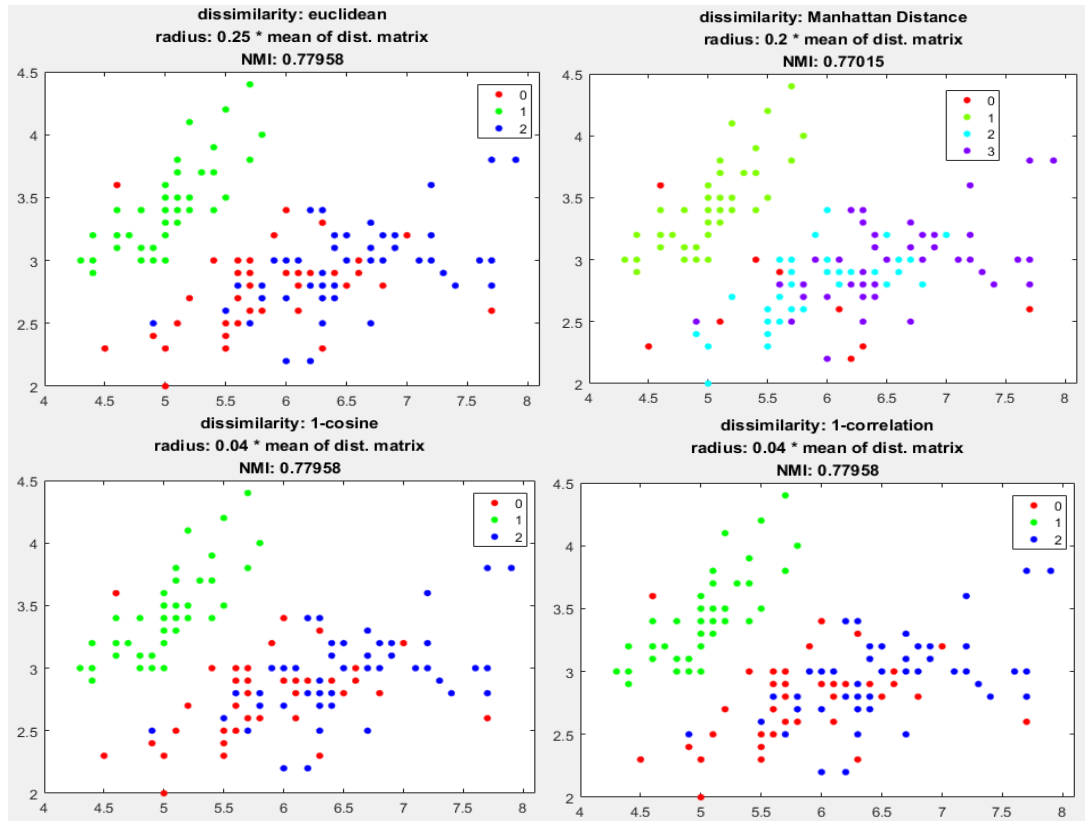


Figure 5. The clustering results for DBSCAN.




Figure 6. The clustering results for DBSCAN with different clustering parameters.

Appendix

```

79 %proximity measures
80 function [DIST]=Data2Relational(X,method)
81
82 %initial output distance array
83 DIST=zeros(length(X(:,1)),length(X(:,1)));
84

```

```

85 if strcmp(method,'euclidean')
86     for i=1:length(X(:,1))
87         for j=1:length(X(:,1))
88             DIST(i,j)=dist(X(i,:),X(j,:));
89         end
90     end
91 end

```

**Euclidean
distance**

```

93 if strcmp(method,'Manhattan')
94     for i=1:length(X(:,1))
95         for j=1:length(X(:,1))
96             for k=1:length(X(1,:))
97                 DIST(i,j)=DIST(i,j)+abs(X(i,k)-X(j,k));
98             end
99         end
100     end
101 end

```

**Manhattan
distance**

```

102
103 if strcmp(method,'cosine')
104     for i=1:length(X(:,1))
105         for j=1:length(X(:,1))
106             numerator=dot(X(i,:),X(j,:));
107             denominator=dist(X(i,:),0)*dist(X(j,:),0);
108             DIST(i,j)=numerator/denominator;
109         end
110     end
111 end

```

**Cosine
similarity**

```

113 if strcmp(method,'correlation')
114     for i=1:length(X(:,1))
115         for j=1:length(X(:,1))
116             co=corrcoef(X(i,:),X(j,:));
117             DIST(i,j)=co(1,2);
118         end
119     end
120 end

```

**Correlation
test**

```

121
122 end

```

```

123 %create the linkage for hierarchical clustering
124 function [Linkage] = HierarchicalClusterLinkage(X,dist,method)
125 %each linkage includes three components
126 %i-th node, j-th node, dist(i-th node, j-th node)
127 %index1,index2, dist(i,j) <= use these variables
128 Linkage=zeros(length(X(:,1)),3);
129
130 %label starts from length(X(:,1))+1 to 2*length(X(:,1))-1
131 %for example, # data points = 3
132 %5 (3 original + 2 merged nodes) nodes will be shown in dendrogram
133 %2 merged node indices are 4 and 5 resp.
134 label=length(X(:,1))+1;
135
136 %record the nodes visited
137 visited=zeros(length(X(:,1)),length(X(:,1)));
138
139 %merge for (m-1) times, where m is # data points
140 for T=1:length(X(:,1))-1
141
142     if strcmp(method,'single')
143         %find smallest element from upper trianglur area of dist
144         MIN=10e10;
145         for i=1:length(X(:,1))
146             for j=i+1:length(X(:,1))
147                 if(dist(i,j)<MIN && visited(i,j)==0)
148                     MIN=dist(i,j);
149                     Index1=i;
150                     Index2=j;
151                 end
152             end
153         end
154
155         %update dist
156         for i=Index1+1:length(X(:,1))
157             if(X(Index2,i)<X(Index1,i) && index2~=i)
158                 X(Index1,i)=X(Index2,i);
159             end
160         end
161         for i=Index2+1:length(X(:,1))
162             if(X(Index1,i)<X(Index2,i) && index1~=i)
163                 X(Index2,i)=X(Index1,i);
164             end
165         end

```

**Single-linkage
clustering**


```

167 -         if strcmp(method,'complete')
168 -             %find smallest element from upper trianglur area of dist
169 -             MIN=10e10;
170 -             for i=1:length(X(:,1))
171 -                 for j=i+1:length(X(:,1))
172 -                     if(dist(i,j)<MIN && visited(i,j)==0)
173 -                         MIN=dist(i,j);
174 -                         Index1=i;
175 -                         Index2=j;
176 -                     end
177 -                 end
178 -             end
179 -             %update dist
180 -             for i=Index1+1:length(X(:,1))
181 -                 if(X(Index2,i)>X(Index1,i) && index2~=i)
182 -                     X(Index1,i)=X(Index2,i);
183 -                 end
184 -             end
185 -             for i=Index2+1:length(X(:,1))
186 -                 if(X(Index1,i)>X(Index2,i) && index1~=i)
187 -                     X(Index2,i)=X(Index1,i);
188 -                 end
189 -             end
190 -         end
191 -
192 -         d=dist(Index1,Index2);
193 -         UseNewLabel=false;
194 -         %if it met the visited nodes (one of one node pair)
195 -         %then using the new linkage label
196 -         for i=1:length(X(:,1))
197 -             if(visited(X(Index1,i))==1)
198 -                 label=label+1;
199 -                 UseNewLabel=true;
200 -                 break;
201 -             end
202 -         end

```

**Complete-linkage
clustering**

```

204 -         for i=1:length(X(:,1))
205 -             if(visited(X(Index2,i))==1)
206 -                 label=label+1;
207 -                 UseNewLabel=true;
208 -                 break;
209 -             end
210 -         end
211
212 -         %the first two components of Linkage are in ascending order
213 -         if(Index1>Index2)
214 -             swap(Index1,Index2);
215 -         end
216
217 -         %put the new label at the second component of Linkage
218 -         if(UseNewLabel)
219 -             Index2=label;
220 -         end
221
222 -         Linkage(T,:)=[Index1, Index2, d];

```

**Linkage
for
dendrogram**

```

223 -     end
224 - end
227 - function [OutputClust]=ClusteringByLinkage(X,Linkage,k)
228
229 -     OutputClust=zeros(length(X(:,1)),length(X(1,:))+1);
230 -     %calculate the height of the tree
231 -     treeHeight=0;
232 -     for i=1:length(Linkage(:,1))
233 -         treeHeight=treeHeight+Linkage(i,3);
234 -     end
235

```

```

236 -     %if we want k clusters, then we can give a cut
237 -     %between k-from-last and (k-1)-from-last nodes
238
239 -     %k-from-last node indices
240 -     [k_Idx1,k_Idx2,~]=Linkage(k,:);
241 -     % (k-1)-from-last node indices
242 -     [k2_Idx1,k2_Idx2,~]=Linkage(k,:);
243
244 -     for i=1:length(X(:,1))
245 -         for j=1:k
246 -             if(traversal(X,k_Idx1) && traversal(X,k_Idx2) && ...
247 -                 traversal(X,k2_Idx1) && traversal(X,k2_Idx2))
248 -                 OutputClust=[X(i,:),j];
249 -                 break;
250 -             end
251 -         end
252 -     end
253 - end

```

**Clustering
by
linkages**

```

254 function [label, noise]=DBSCAN(X, radius, Pthres, method)
255
256 %{
257     radius: maximum distance between two samples for one to
258           be considered as in the neighborhood of the other
259     Pthres: # samples (including the point itself) in one neighborhood
260           for a point to be considered as a core point. This
261
262     method: Euclidean, Manhattan, cosine, correlation
263     %}
264
265     %idx for clustering
266     %idx=0 => noise
267     idx=0;
268
269     %initial output clustering label
270     label=zeros(length(X(:,1)),1);
271     D=pdist2(X,X,method);
272     %visit or not
273     visited=false(n,1);
274     %is noise or not
275     noise=false(n,1);
276
277     %traversal all data points
278     for i=1:n
279         if ~visited(i)
280             visited(i)=true;
281
282             %check if # of neighbors < Pthres
283             %if so, then it will be labeled as "noise"
284             if (numel(FindNeighbor(i))<Pthres)
285                 noise(i)=true;
286             else
287                 idx=idx+1;
288                 %find iteratively
289                 FindMore(i,Nb,idx);
290             end
291         end
292     end
293 end

```

DBSCAN
Find neighbors

```

295 %now we want to see more data points
296 function FindMore(i,Neighbors,C)
297     label(i)=C;
298     k=1;
299     while true
300         %traversal every data points' neighbors in the current radius
301         j = Neighbors(k);
302         if ~visited(j)
303             visited(j)=true;
304             NeighborsNeighbor=FindNeighbor(j);
305             if numel(NeighborsNeighbor)>=Pthres
306                 %bind all the data points we find together
307                 Neighbors=[Neighbors NeighborsNeighbor]; %#ok
308             end
309         end
310         %cluster these data points into same group
311         if (label(j)==0)
312             label(j)=C;
313         end
314         k=k+1;
315         %if traversal all the neighbors
316         if (k>numel(Neighbors))
317             break;
318         end
319     end
320 end
321
322 function Nb=FindNeighbor(i)
323     Nb=find(D(i,:)<=radius);
324 end
325
326 end

```

DBSCAN
Find neighbors'
neighbors